

Machine Learning

Professor: Dr. Raman Kannan



Homework 2 : Ensemble Learning

Topic : Music Genre Prediction

Aakash Shetty
aks9108



Outline

1. Overview
2. Preprocessing Recap
3. Train & Test Split
4. Random Forest
5. Boosting
6. Bagging
7. Comparison
8. What does cross validation do to bias and variance?
9. Conclusion



Overview

Let's start with data, we have a 18 attribute columns, of which one is the target column 'music_genre'. In this homework we will be going through three ensemble techniques. Random Forest, Boosting and Bagging.

Given these 3 techniques we will try to compare them with each other based on accuracy, variance and bias. Variance is higher for over-fitting and bias will be higher for under-fitting.

Based on what we know random forest and bagging technique will try to reduce variance while boosting will try to reduce the bias. Let's see if we can see the difference in actual data we have worked with.

Preprocessing Recap

- The data used is cleaned by removing NULL values and correcting the schema incase there are data type mismatch. We also remove the outliers that may affect the training.
- Next, we made some changes to the data by encoding the key and mode attribute columns. We change the character values to numeric values for further processing. 'Major' and 'Minor' was mapped to 1 and 0 respectively in Mode while, in keys 'A', 'B', 'C#' etc. was mapped to values from 1-12.
- We also looked at important attributes for training testing in EDA which can be useful in for modeling. We disregard such columns and consider only columns which will be used for modeling. Columns such as song name, artist, popularity and instance id will not be considered for modeling.
- Important features include energy, danceability, instumentalness, loudness etc. In total we have 11 attributes to be considered for predicting a song out of 10 genres.
- We will be using the data from EDA as the input data for train and test spilt.

```
songs <- read.csv('processed_songs.csv', stringsAsFactors = TRUE)
head(songs, 10)
```

A data.frame: 10 x 18

	instance_id	artist_name	track_name	popularity	acousticness	danceability	duration_ms	energy	instrumentalness	key	liveness	loudness	mode	speech
	<int>	<fct>	<fct>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<dbl>	<dbl>	<int>	
1	32894	Röyksopp	Röyksopp's Night Out	0	0.00468	0.652	-1	0.941	7.92e-01	2	0.1150	-5.201	1	
2	46652	Thievery Corporation	The Shining Path	0	0.01270	0.622	218293	0.890	9.50e-01	6	0.1240	-7.043	1	
3	30097	Dillon Francis	Hurricane	0	0.00306	0.620	215613	0.755	1.18e-02	12	0.5340	-4.617	0	
4	62177	Dubloadz	Nitro	0	0.02540	0.774	166875	0.700	2.53e-03	5	0.1570	-4.498	0	
5	24907	What So Not	Divide & Conquer	0	0.00465	0.638	222369	0.587	9.09e-01	10	0.1570	-6.266	0	
6	43760	Jordan Comolli	Clash	0	0.02890	0.572	214408	0.803	7.74e-06	3	0.1060	-4.294	0	
7	30738	Hraach	Delirio	0	0.02970	0.809	416132	0.706	9.03e-01	11	0.0635	-9.339	1	
8	84950	Kayzo	NEVER ALONE	0	0.00299	0.509	292800	0.921	2.76e-04	9	0.1780	-3.175	1	
9	56950	Shlump	Lazer Beam	0	0.00934	0.578	204800	0.731	1.12e-02	1	0.1110	-7.091	1	

Train & Test Split

We split the train data and test data into 75:25 ratio. All the features will be taken as numeric values.

Split Train and Test Data

```
set.seed(11111)
feats <- names(songs)[c(5:11,13:15,17)]
train_songs <- songs %>%
  mutate_if(is.numeric, scale)

training_songs <- sample(1:nrow(train_songs), nrow(train_songs)*.75, replace = FALSE)
train_set <- train_songs[training_songs, c('music_genre', feats)]
test_set <- train_songs[-training_songs, c('music_genre', feats)]
```

feats

'acousticness' · 'danceability' · 'duration_ms' · 'energy' · 'instrumentalness' · 'key' · 'liveness' · 'mode' · 'speechiness' · 'tempo' · 'valence'

str(train_set)

```
'data.frame':  33496 obs. of  12 variables:
 $ music_genre      : Factor w/ 10 levels "Alternative",...: 9 4 3 10 2 3 7 1 8 1 ...
 $ acousticness     : num [1:33496, 1] -0.755 1.925 1.845 -0.794 -0.888 ...
 $ danceability     : num [1:33496, 1] 0.881 -0.444 0.447 -0.601 -1.052 ...
 $ duration_ms      : num [1:33496, 1] -2.0309 0.7316 -0.8379 0.0961 -2.0309 ...
 $ energy           : num [1:33496, 1] 0.345 -1.443 -2.071 1.137 1.027 ...
 $ instrumentalness : num [1:33496, 1] -0.552 -0.417 -0.305 -0.366 -0.552 ...
 $ key             : num [1:33496, 1] -0.968 -1.256 0.476 -1.545 -0.101 ...
 $ liveness         : num [1:33496, 1] -0.3902 -0.6618 -0.5018 0.0254 -0.6023 ...
 $ mode            : num [1:33496, 1] -0.748 1.337 -0.748 -0.748 -0.748 ...
 $ speechiness      : num [1:33496, 1] -0.592 -0.531 -0.544 -0.58 -0.588 ...
 $ tempo           : num [1:33496, 1] 0.324 1.21 0.682 0.147 0.813 ...
 $ valence          : num [1:33496, 1] 0.156 -0.6638 -1.3497 1.337 -0.0754 ...
```

str(test_set)

```
'data.frame':  11166 obs. of  12 variables:
 $ music_genre      : Factor w/ 10 levels "Alternative",...: 6 6 6 6 6 6 6 6 6 ...
 $ acousticness     : num [1:11166, 1] -0.809 -0.847 1.642 -0.894 -0.811 ...
 $ danceability     : num [1:11166, 1] 0.0636 0.5596 0.9936 -1.7005 0.5371 ...
 $ duration_ms      : num [1:11166, 1] -0.0111 0.1931 -2.0309 -0.285 0.8642 ...
 $ energy           : num [1:11166, 1] 0.764 1.103 -0.75 0.323 0.193 ...
 $ instrumentalness : num [1:11166, 1] -0.551 -0.55 -0.44 -0.525 2.1 ...
 $ key             : num [1:11166, 1] -0.968 -0.679 -1.545 1.631 1.053 ...
 $ liveness         : num [1:11166, 1] -0.545 -0.192 -0.13 0.894 -0.508 ...
 $ mode            : num [1:11166, 1] -0.748 -0.748 1.337 -0.748 1.337 ...
 $ speechiness      : num [1:11166, 1] 2.524 -0.291 -0.508 -0.519 -0.536 ...
 $ tempo           : num [1:11166, 1] 0.9769 0.2603 1.1319 0.5048 -0.0676 ...
 $ valence          : num [1:11166, 1] -0.928 -1.411 0.765 -1.398 -1.053 ...
```

Random Forest

Random forest is an ensemble model using bagging as the ensemble method and decision tree as the individual model. One random subset is used to train one decision tree. The optimal splits for each decision tree are based on a random subset of features. Each individual tree predicts the records in the test set, independently.

MODEL

```
songs_rf <- randomForest(music_genre~., data = train_set, mtry = 4)

pred_train <- predict(songs_rf)
pred_test  <- predict(songs_rf, test_set)
```

TRAIN

```
confusionMatrix(pred_train, as.factor(train_set$music_genre))
```

Confusion Matrix and Statistics

	Reference							
Prediction	Alternative	Anime	Blues	Classical	Country	Electronic	Hip-Hop	Jazz
Alternative	658	304	233	94	145	198	146	87
Anime	348	1577	207	156	159	183	30	125
Blues	292	178	1344	50	323	157	59	479
Classical	9	381	39	2580	5	12	3	241
Country	343	264	422	11	1878	63	75	122
Electronic	280	223	151	55	32	2062	130	435
Hip-Hop	328	31	85	4	136	152	1335	137
Jazz	204	200	471	169	126	370	61	1625
Rap	304	38	53	0	136	86	1526	82
Rock	559	184	363	19	427	97	33	55

	Reference	
Prediction	Rap	Rock
Alternative	243	636
Anime	65	298
Blues	37	606
Classical	0	27
Country	124	827
Electronic	120	179
Hip-Hop	1914	81
Jazz	58	189
Rap	691	119
Rock	103	435

TRAIN

Overall Statistics	
Accuracy	: 0.4235
95% CI	: (0.4182, 0.4288)
No Information Rate	: 0.1014
P-Value [Acc > NIR]	: < 2.2e-16
Kappa	: 0.3594
McNemar's Test P-Value	: NA

TRAIN

Statistics by Class:				
	Class: Alternative	Class: Anime	Class: Blues	
Sensitivity	0.19789	0.46657	0.39905	
Specificity	0.93086	0.94784	0.92761	
Pos Pred Value	0.23980	0.50095	0.38128	
Neg Pred Value	0.91327	0.94059	0.93247	
Prevalence	0.09927	0.10091	0.10055	
Detection Rate	0.01964	0.04708	0.04012	
Detection Prevalence	0.08192	0.09398	0.10524	
Balanced Accuracy	0.56438	0.70720	0.66333	
	Class: Classical	Class: Country	Class: Electronic	
Sensitivity	0.82218	0.55777	0.61006	
Specificity	0.97638	0.92529	0.94671	
Pos Pred Value	0.78253	0.45483	0.56231	
Neg Pred Value	0.98152	0.94930	0.95581	
Prevalence	0.09368	0.10052	0.10091	
Detection Rate	0.07702	0.05607	0.06156	
Detection Prevalence	0.09843	0.12327	0.10948	
Balanced Accuracy	0.89928	0.74153	0.77838	
	Class: Hip-Hop	Class: Jazz	Class: Rap	Class: Rock
Sensitivity	0.39288	0.47963	0.20596	0.12805
Specificity	0.90471	0.93862	0.92223	0.93887
Pos Pred Value	0.31763	0.46790	0.22768	0.19121
Neg Pred Value	0.92957	0.94128	0.91254	0.90513
Prevalence	0.10144	0.10115	0.10016	0.10142
Detection Rate	0.03986	0.04851	0.02063	0.01299
Detection Prevalence	0.12548	0.10368	0.09061	0.06792
Balanced Accuracy	0.64879	0.70913	0.56410	0.53346

TEST

confusionMatrix(pred_test, as.factor(test_set\$music_genre))								
Confusion Matrix and Statistics								
	Reference							
Prediction	Alternative	Anime	Blues	Classical	Country	Electronic	Hip-Hop	Jazz
Alternative	247	100	60	33	50	62	60	33
Anime	114	546	66	65	48	64	9	45
Blues	87	62	410	15	111	48	14	159
Classical	4	97	17	912	3	3	0	85
Country	122	82	115	7	626	22	19	46
Electronic	102	83	44	11	7	635	37	129
Hip-Hop	131	12	33	0	48	49	453	55
Jazz	66	79	185	63	28	118	23	514
Rap	106	8	14	0	52	29	496	18
Rock	189	39	121	5	146	38	10	15
	Reference							
Prediction	Rap	Rock						
Alternative	70	249						
Anime	19	99						
Blues	10	195						
Classical	1	11						
Country	44	262						
Electronic	37	70						
Hip-Hop	682	28						
Jazz	18	61						
Rap	231	37						
Rock	37	146						

TEST

Overall Statistics	
Accuracy : 0.4227	
95% CI : (0.4135, 0.4319)	
No Information Rate : 0.1046	
P-Value [Acc > NIR] : < 2.2e-16	
Kappa : 0.3587	
McNemar's Test P-Value : NA	

TEST				
Statistics by Class:				
	Class: Alternative	Class: Anime	Class: Blues	
Sensitivity	0.21147	0.49278	0.38498	
Specificity	0.92829	0.94741	0.93060	
Pos Pred Value	0.25622	0.50791	0.36904	
Neg Pred Value	0.90972	0.94431	0.93486	
Prevalence	0.10460	0.09923	0.09538	
Detection Rate	0.02212	0.04890	0.03672	
Detection Prevalence	0.08633	0.09627	0.09950	
Balanced Accuracy	0.56988	0.72009	0.65779	
	Class: Classical	Class: Country	Class: Electronic	
Sensitivity	0.82088	0.55943	0.59457	
Specificity	0.97802	0.92844	0.94850	
Pos Pred Value	0.80494	0.46543	0.54978	
Neg Pred Value	0.98017	0.94980	0.95675	
Prevalence	0.09950	0.10021	0.09565	
Detection Rate	0.08168	0.05606	0.05687	
Detection Prevalence	0.10147	0.12045	0.10344	
Balanced Accuracy	0.89945	0.74393	0.77154	
	Class: Hip-Hop	Class: Jazz	Class: Rap	Class: Rock
Sensitivity	0.40410	0.46770	0.20104	0.12608
Specificity	0.89667	0.93633	0.92413	0.94005
Pos Pred Value	0.30382	0.44502	0.23310	0.19571
Neg Pred Value	0.93096	0.94156	0.90978	0.90288
Prevalence	0.10039	0.09842	0.10290	0.10371
Detection Rate	0.04057	0.04603	0.02069	0.01308
Detection Prevalence	0.13353	0.10344	0.08875	0.06681
Balanced Accuracy	0.65038	0.70201	0.56259	0.53306

VARIANCE & BIAS
<code>var(as.numeric(pred_test), as.numeric(test_set\$music_genre))</code>
1.63257517001703
<code>bias(as.numeric(pred_test), as.numeric(test_set\$music_genre))</code>
-0.0802435966326348

Observation:

Accuracy	42%
Variance	1.632
Bias	-0.08

- We get a low accuracy of 42% with reasonable variance and bias.
- The data still continues to show low accuracy with models which goes to show that the data would need a more complex model to get good results like a neural network or transformer.

Boosting

Boosting is an ensemble learning method that combines a set of weak learners into a strong learner to minimize training errors. In boosting, a random sample of data is selected, fitted with a model and then trained sequentially, each model tries to compensate for the weaknesses of its predecessor. We have used XGBoost.

MODEL

```
matrix_train_gb <- xgb.DMatrix(data = as.matrix(train_set[,-1]), label = as.integer(as.factor(train_set[,1])))
matrix_test_gb <- xgb.DMatrix(data = as.matrix(test_set[,-1]), label = as.integer(as.factor(test_set[,1])))

model_gb <- xgboost(data = matrix_train_gb,
                    nrounds = 50,
                    verbose = FALSE,
                    params = list(objective = "multi:softmax",
                                num_class = 10 + 1))

predict_gb_one <- predict(model_gb, matrix_test_gb)
predict_gb <- levels(as.factor(test_set$music_genre))[predict_gb_one]
```

TEST

confusionMatrix(as.factor(predict_gb), as.factor(test_set\$music_genre))

Confusion Matrix and Statistics

	Reference								
Prediction	Alternative	Anime	Blues	Classical	Country	Electronic	Hip-Hop	Jazz	
Alternative	218	96	68	25	40	59	30	20	
Anime	124	543	55	75	37	63	9	48	
Blues	92	51	432	15	124	48	15	168	
Classical	12	102	19	914	5	6	1	83	
Country	136	89	125	6	641	23	19	41	
Electronic	115	79	37	11	9	632	33	118	
Hip-Hop	118	4	28	0	48	54	509	48	
Jazz	79	74	175	56	27	104	28	526	
Rap	100	14	19	0	50	35	462	23	
Rock	174	56	107	9	138	44	15	24	

	Reference	
Prediction	Rap	Rock
Alternative	68	168
Anime	21	120
Blues	14	211
Classical	2	10
Country	42	265
Electronic	33	78
Hip-Hop	631	28
Jazz	24	67
Rap	286	31
Rock	28	180

TEST	
Overall Statistics	
Accuracy : 0.4371	
95% CI : (0.4279, 0.4464)	
No Information Rate : 0.1046	
P-Value [Acc > NIR] : < 2.2e-16	
Kappa : 0.3748	
McNemar's Test P-Value : < 2.2e-16	

TEST				
Statistics by Class:				
Class: Alternative Class: Anime Class: Blues				
Sensitivity	0.18664	0.49007	0.40563	
Specificity	0.94259	0.94512	0.92694	
Pos Pred Value	0.27525	0.49589	0.36923	
Neg Pred Value	0.90842	0.94390	0.93667	
Prevalence	0.10460	0.09923	0.09538	
Detection Rate	0.01952	0.04863	0.03869	
Detection Prevalence	0.07093	0.09807	0.10478	
Balanced Accuracy	0.56462	0.71760	0.66629	
Class: Classical Class: Country Class: Electronic				
Sensitivity	0.82268	0.57283	0.59176	
Specificity	0.97613	0.92575	0.94920	
Pos Pred Value	0.79203	0.46215	0.55197	
Neg Pred Value	0.98032	0.95112	0.95649	
Prevalence	0.09950	0.10021	0.09565	
Detection Rate	0.08186	0.05741	0.05660	
Detection Prevalence	0.10335	0.12422	0.10254	
Balanced Accuracy	0.89941	0.74929	0.77048	
Class: Hip-Hop Class: Jazz Class: Rap Class: Rock				
Sensitivity	0.45406	0.47862	0.24891	0.15544
Specificity	0.90453	0.93702	0.92672	0.94055
Pos Pred Value	0.34673	0.45345	0.28039	0.23226
Neg Pred Value	0.93689	0.94273	0.91494	0.90588
Prevalence	0.10039	0.09842	0.10290	0.10371
Detection Rate	0.04558	0.04711	0.02561	0.01612
Detection Prevalence	0.13147	0.10389	0.09135	0.06941
Balanced Accuracy	0.67929	0.70782	0.58782	0.54799

VARIANCE & BIAS
var(as.numeric(predict_gb_one), as.numeric(test_set\$music_genre))
1.85823707196347
bias(as.numeric(predict_gb_one), as.numeric(test_set\$music_genre))
-0.0167472684936414

Observation:

Accuracy	44%
Variance	1.86
Bias	-0.016

- We get a low accuracy of 43% with reasonable variance and negative bias.
- The data still continues to show low accuracy with models.
- Variance is higher than random forest , showing a more overfitting.
- XGBoost does not need normalized features and work well if the data is nonlinear, non-monotonic, or with segregated clusters.

Bagging

Bagging is an ensemble learning method that is commonly used to reduce variance within a noisy dataset. In bagging, a random sample of data in a training set is selected with replacement, meaning that the individual data points can be chosen more than once.

MODEL

```
gbag <- bagging(music_genre ~ ., data = train_set, coob=TRUE)
predict_bag <- predict(gbag, newdata=test_set)
```

TEST

```
confusionMatrix(as.factor(predict_bag), as.factor(test_set$music_genre))
```

Confusion Matrix and Statistics

	Reference							
Prediction	Alternative	Anime	Blues	Classical	Country	Electronic	Hip-Hop	Jazz
Alternative	214	100	67	28	71	78	66	37
Anime	116	516	75	87	51	58	11	40
Blues	85	59	363	23	122	55	16	157
Classical	9	91	14	886	2	2	1	82
Country	117	88	120	6	561	24	15	41
Electronic	100	83	50	11	11	592	44	137
Hip-Hop	133	17	31	0	50	49	390	52
Jazz	56	76	175	61	24	138	22	483
Rap	99	7	20	2	55	29	542	31
Rock	239	71	150	7	172	43	14	39

	Reference	
Prediction	Rap	Rock
Alternative	70	246
Anime	25	123
Blues	15	183
Classical	0	5
Country	41	240
Electronic	37	70
Hip-Hop	649	27
Jazz	20	56
Rap	246	43
Rock	46	165

TEST

Overall Statistics	
Accuracy	: 0.3955
95% CI	: (0.3864, 0.4046)
No Information Rate	: 0.1046
P-Value [Acc > NIR]	: < 2.2e-16
Kappa	: 0.3284
Mcnemar's Test P-Value	: < 2.2e-16

TEST				
Statistics by Class:				
	Class: Alternative	Class: Anime	Class: Blues	
Sensitivity	0.18322	0.46570	0.34085	
Specificity	0.92368	0.94174	0.92921	
Pos Pred Value	0.21904	0.46824	0.33673	
Neg Pred Value	0.90637	0.94118	0.93041	
Prevalence	0.10460	0.09923	0.09538	
Detection Rate	0.01917	0.04621	0.03251	
Detection Prevalence	0.08750	0.09869	0.09654	
Balanced Accuracy	0.55345	0.70372	0.63503	
	Class: Classical	Class: Country	Class: Electronic	
Sensitivity	0.79748	0.50134	0.55431	
Specificity	0.97951	0.93112	0.94623	
Pos Pred Value	0.81136	0.44773	0.52159	
Neg Pred Value	0.97767	0.94371	0.95255	
Prevalence	0.09950	0.10021	0.09565	
Detection Rate	0.07935	0.05024	0.05302	
Detection Prevalence	0.09780	0.11222	0.10165	
Balanced Accuracy	0.88850	0.71623	0.75027	
	Class: Hip-Hop	Class: Jazz	Class: Rap	Class: Rock
Sensitivity	0.34790	0.43949	0.21410	0.14249
Specificity	0.89965	0.93762	0.91734	0.92196
Pos Pred Value	0.27897	0.43474	0.22905	0.17442
Neg Pred Value	0.92516	0.93874	0.91052	0.90284
Prevalence	0.10039	0.09842	0.10290	0.10371
Detection Rate	0.03493	0.04326	0.02203	0.01478
Detection Prevalence	0.12520	0.09950	0.09618	0.08472
Balanced Accuracy	0.62378	0.68855	0.56572	0.53222

VARIANCE & BIAS
<pre>var(as.numeric(predict_bag), as.numeric(test_set\$music_genre))</pre>
1.44233802971226
<pre>bias(as.numeric(predict_bag), as.numeric(test_set\$music_genre))</pre>
0.00644814615797958

Observation:

Accuracy	39%
Variance	1.4
Bias	0.006

- We get a low accuracy of 39% with lower variance and higher bias compared to the other two models.
- This can be because the model tries to generalize the data and but it fails to train well enough with given techniques.

Comparison

<i>Model</i>	<i>Accuracy</i>	<i>Variance</i>	<i>Bias</i>
Random Forest	42%	1.632	-0.08
Boosting	44%	1.86	-0.016
Bagging	39%	1.4	0.006

- XGBoost has the best accuracy but still fails to give promising results for the dataset.
- Bagging will try to decrease variance as you can see bagging technique has lower variance.
- Boosting will try to reduce the Bias mainly thus giving us lower bias than other bagging techniques.

What does cross validation do to bias and variance?

Cross Validation techniques reduce over-fitting, it reduces the variance while also trying to reduce the bias. However, we know about the tradeoff between variance and bias. As the bias increase, the variance reduce and vis-a-versa.

Conclusion

As discussed in the overview we can see the difference in the results for the given techniques. While bagging techniques have lower variance and higher bias, boosting technique has higher bias and lower variance in comparison. Ensemble techniques seem to work better than many models implemented in HW1.